

Multicore Programming
Project Erlang
Key-value store

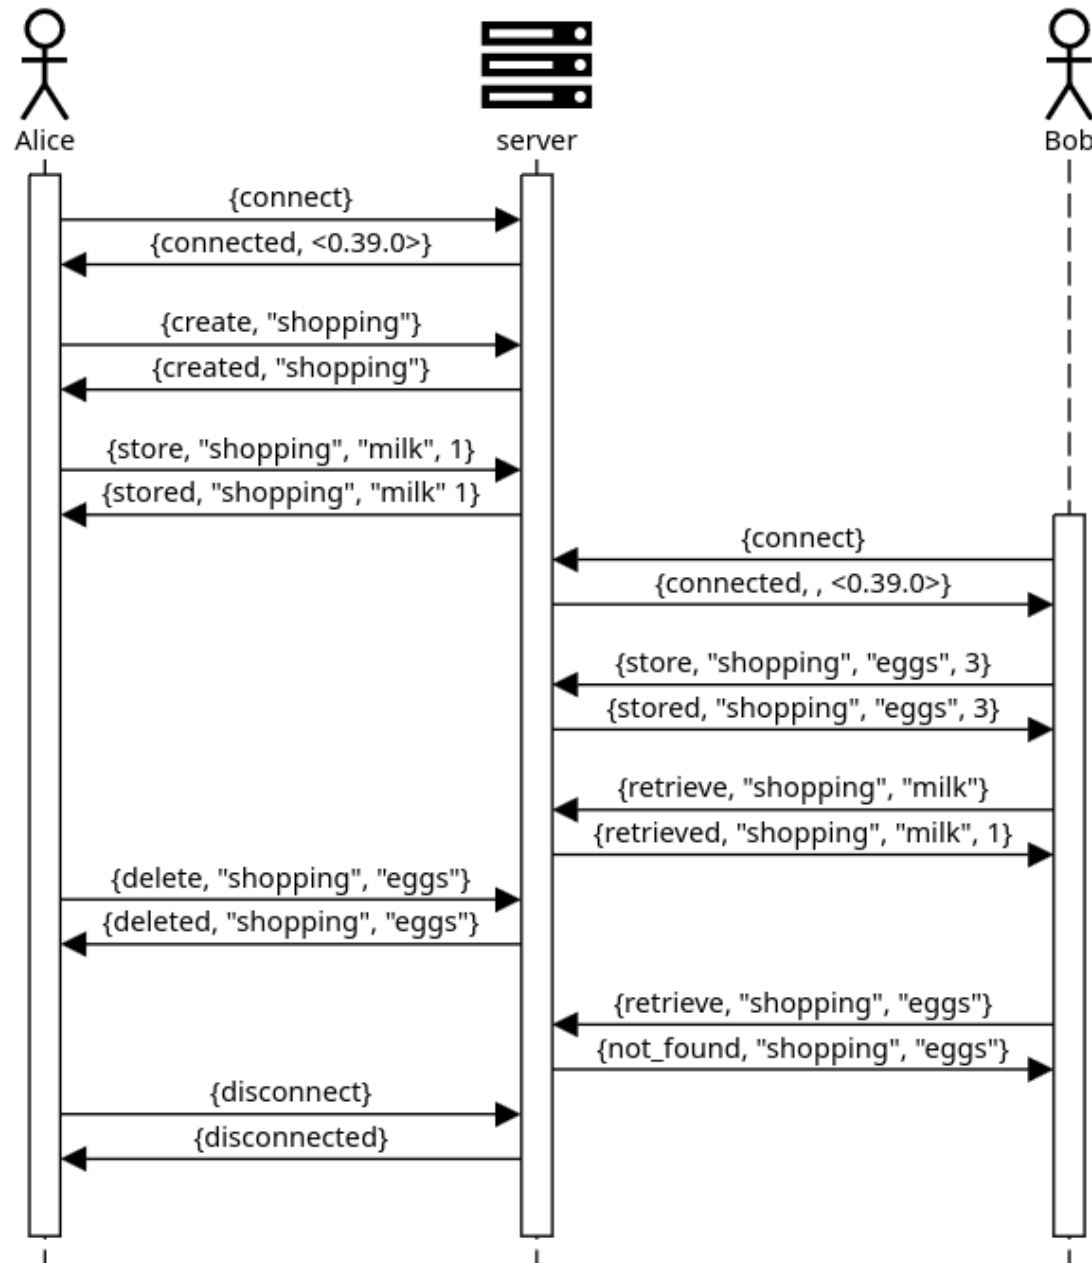
Key-value store

A key-value store is a database in which you associate a key with a value.

You will need to create:

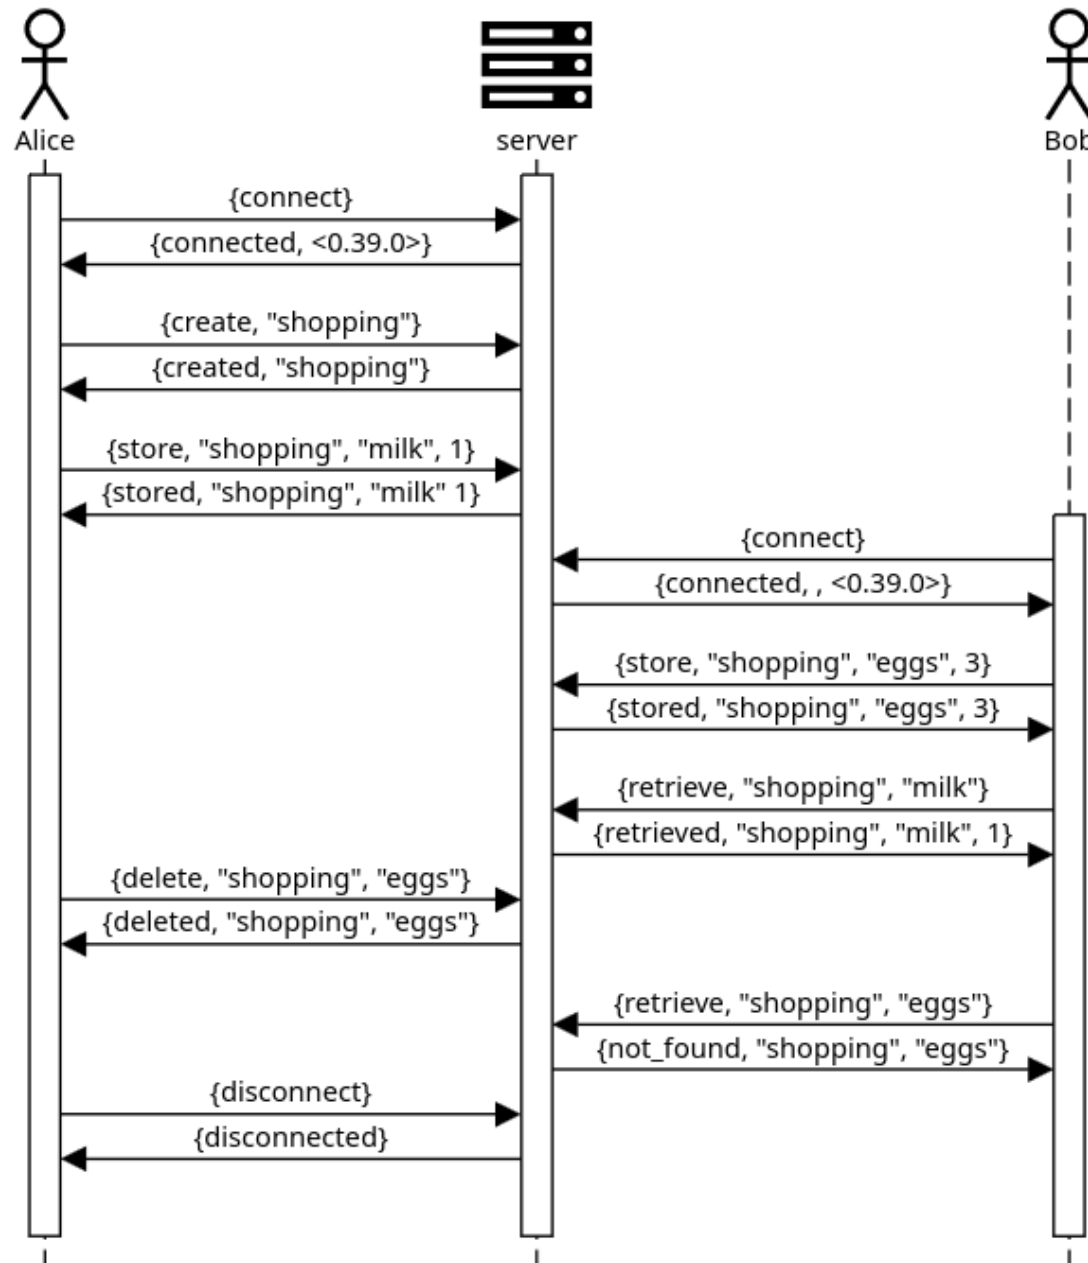
- ◆ Implementation
- ◆ Evaluation
- ◆ Report

Example session



Key-value pairs
are stored in
buckets.

Example session




Operations:

- connect
- disconnect
- create
- store
- retrieve
 - can return not_found

Example implementation

An example implementation is given with one server.
You need to extend this to have multiple server processes.



- ◆ 1 process / client → don't change, for benchmarking
- ◆ 1 server process → you'll need to change this

Goal: increase scalability

Your goal is to extend the implementation to improve scalability/performance, by using multiple processes.

Possible directions:

- **Sharding** (splitting data into subsets, divided over different processes)
 - Static vs. dynamic distribution?
- **Replication** (copying of data over processes)
 - Strong vs. eventual consistency?
- **Load balancing** incoming requests from clients
 - Static vs. dynamic?
- **Caching**
 - How to handle invalidation?
- ... You are free to explore other directions or combine the above techniques.

Evaluation

Benchmarks: generate many client processes that each do a number of operations (store, retrieve).

You can measure metrics like latency & throughput of operations, in variety of situations.

You should do three experiments:

- Speed-up for increasing number of threads
- “Best-case” scenario
- A scenario you can choose freely
(ideas: change types of requests, change R/W ratio, change # buckets, change # keys/bucket...)

We provide a basic benchmark in Erlang and a matplotlib script in Python to process the results.

Evaluation: hardware

Run experiments on:

- your machine or machine in computer room
(requirement: ≥ 4 cores)
- “Firefly”: 64-core server at lab
(we will arrange remote access)

Report

Table of contents in assignment sheet:

Implementation

Describe architecture (incl. diagrams)

How do you ensure scalability?

Evaluation

Describe set-up, metrics

Include plots of results

Explain what you see and why you see it

Insight questions

Hypothetical extensions

Details

Deadline: Thursday 10th of April, 23:59

Submit ZIP with implementation & report on Canvas

Project defense in June

$\frac{1}{3}$ of your final grade

Per day too late, -2 points

We check for plagiarism

Details on Canvas