KEY-VALUE STORE

Multicore Programming: Project Erlang

Gérard Lichtert
gerard.lichtert@vub.be

March 31, 2025

Sciences and Bio-Engineering Sciences

Contents

1   Overview

To give a brief summary of the reports' contents, we will first start with the implementation, which discusses a sharded and cached implementation (henceforth called the 'sharded' implementation). The sharded implementation distributes each store in an actor (called bucket) and keeps a central actor to coordinate communication between clients and the buckets. Another implementation that is discussed is a worker, sharded and cached implementation (henceforth called the 'worker' implementation). This implementation is similar to the previous in regards of the buckets, however whenever a client connects to the server, they get their own process to communicate with. So this implementation has a one to one process relation whilst the previous has many to one relation.

These implementations together with the provided central implementation are used to perform some experiments. The first experiment measures the speedup, depending on the number of scheduler threads. The second measures throughput depending on the number of scheduler threads. Both of these experiments use the worker implementation to analyze the results, however the results for the other implementations are also available in the dataset. The last compares the throughput of the different implementations

2   Implementation

2.1   Central implementation

2.1.1   Architecture

To start explaining the implementations, we will start out with the provided central implementation, which is just a single process containing a dictionary of dictionary, or a dictionary of buckets, to which the user sends CRUD (Create, Read, Update, Delete) messages. The single process or 'Server' handles all the messages and also responds to them after performing the necessary operations. A diagram showcasing the flow of the messages can be seen in Figure 1. As this implementation was meant to be improved upon for scalability, we will not be answering the scalability questions for this implementation.

Figure 1: Message flow of the Central implementation

2.2 Sharded-Cached implementation

2.2.1 Architecture

To split the central implementation up, instead of having a nested dictionary, we create a process per bucket, so each bucket process will hold a dictionary or the Key-Value store and the central process will keep track of the process ID's in its dictionary to address each bucket. This means that while the central process still has to handle all communication from the clients, it no longer has to manage the data by itself only the addresses of the buckets. The only responsability that the server has is to create the buckets when needed and forward the messages of the client to the appropiate bucket. The responsability of the buckets is to perform the operations that the server forwards to it and to respond to the clients. A diagram showing the flow of communication can be seen in Figure 2.


Figure 2: Message flow of the Sharded implementation

To improve on the performance of the implementation we also add a caching mechanism that remembers the last bucket addressed in the server. If a message is to be forwarded to the same bucket twice or more in a row, the server no longer needs to look the bucket up since it already has the address. This mechanism is also applied in the buckets for retrieval operations, where it

3

remembers the last value associated with the last lookup.

2.2.2 Scalability

This implementation allows the system to scale with the amount of key-value stores that exist, since each gets its own process to live in, and the load gets split to the buckets themselves. The best case would be if the clients operate on different buckets and the worst case would be if all clients operate on the same bucket. However since the bottleneck is the server, the load on the buckets will at most be the same as the server.

2.3 Worker-Sharded-Cached implementation

To improve on the Sharded implementation, we need to solve the bottleneck in the server. Currenty all communication of all the clients go through the server and get forwarded to the appropiate buckets which is cumbersome for the server. To aleviate the server we create a worker process for each client that connects to the server, and in doing so, each client will have its own worker to which it sends messages to.
The worker will in turn forward the messages to the appropiate bucket. The only issue with this implementation is that when a client creates a bucket, the bucket must also be available to all the other clients connected to the service. To solve this we add the responsability of replicating the address of the bucket to the server. The server keeps track of all existing worker processes and whenever one of them send a replication message (when a new bucket is created), it broadcasts it to the rest of the existing worker processes, which add the bucket to their dictionary of buckets. An example of this is shown in the diagram in Figure 3.
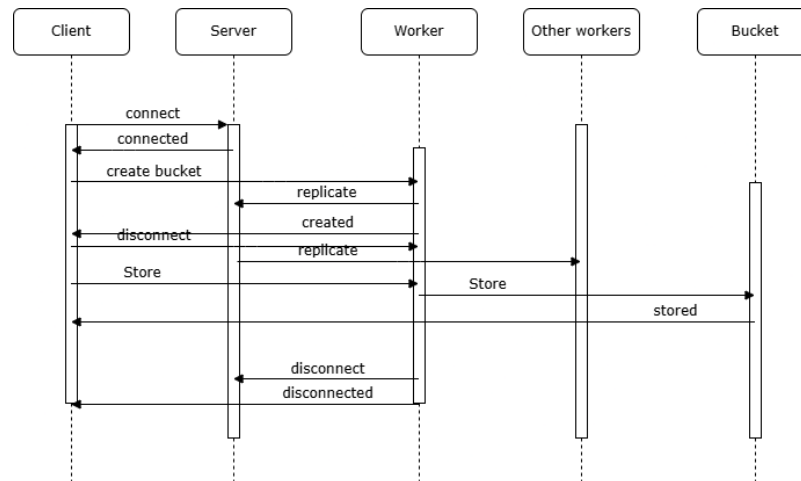

Figure 3: Creating a bucket in the Worker implementation

4

To keep making use of the caching mechanism, it is moved from the server to the worker process, since the workers keep track of the buckets that the client interacts with. Moreover, since it will always be the same client interacting with the worker, the cache will be hit more frequently. A general flow of messages can be seen in Figure 4.


Figure 4: Message flow of the Worker implementation

2.4   Scalability

This implementation is a lot more scalable than the previous implementation. It now allows the system to also scale with the amount of clients and thus allows for messages to be processed in parallel. The best case would be if the client on different buckets, allowing the load to be distributed across the buckets. The worst case would be if all the clients operate on a single bucket, creating a bottleneck in the bucket. Since this implementation allows for the parallelisation of the sent messages, the bottleneck there is now the speed of the processing of the messages. Another bottleneck would be the buckets themselves, if they receive more messages than they can handle. For this we could explore further sharding of the bucket, or replication as well as pla:!cing the bucket behind a load balancer.

5

3 Evaluation

3.1 Experimental set-up

For the experiments I used two devices. Firstly the Firefly machine of the VUB and my own laptop henceforth called "Omen" to set up the experiments for Firefly as well as make the necessary scripts and programs to run the experiments on firfly and analyze the data automatically. The specifications of both devices are found in Table 1

Device	Firefly	Omen HP laptop
Hardware		
CPU	AMD Ryzen Threadripper 3990X Processor (64 cores / 128 threads @ 2.9 GHz Base, 4.3 GHz boost)	Intel Core i7-9750H (6 cores / 12 threads) @ 2.6 GHz, 4.5 Ghz boost
RAM	128 GB (DDR4 3200 MHz)	16GB (DDR4 2667 MHz)
Software		
OS	Ubuntu 24.04.1	Microsoft Windows 11 Home
Erlang/OTP version	25.3.2.8	27

Table 1: Hardware used for the experiments

For the experiments I used a Powershell script and a Bash script to start an erlang process to perform an experiment and benchmark it 30 times. This was done for each implementation, so for the central, sharded and worker implementation and for three different scenarios. The scenarios were: A read only scenario, where clients only send retrieve messages to the server. A read write scenario where clients performed 50-50 writes and reads. Finally also a mixed scenario where the balance was 90% reads and 10% writes. Note that only the time elapsed for performing these operations was measured and thus the buckets were already existing prior to starting the benchmark.

Each experiment has 100 clients performing operations on the system. Each client will do 1000 operations per scenario and 2000 operations in the read-write scenario. There will always be 1000 buckets, each holding 100 keys.

3.2 Experimental methodology

To measure the time elapsed we used the wall clock time on the client side and we performed the experiment 64 times per scenario and per implementation on the Omen device, once per amount of scheduler threads (so 1-64). On Firefly we performed the experiment 128 times per scenario, per implementation since it can use a lot more threads.

3.3 Experiments

For the first experiment we will be comparing the throughput of the different implementations on the maximum amount of scheduler threads (12 for Omen and 128 for Firefly).

3.3.1 Throughput per second depending on the implementation and scenario

For this experiment we use the parameters mentioned in subsection 3.1. For each device we use the measurements for the maximum amount of threads, so in Omen's case its 12 and for Firefly its 128.
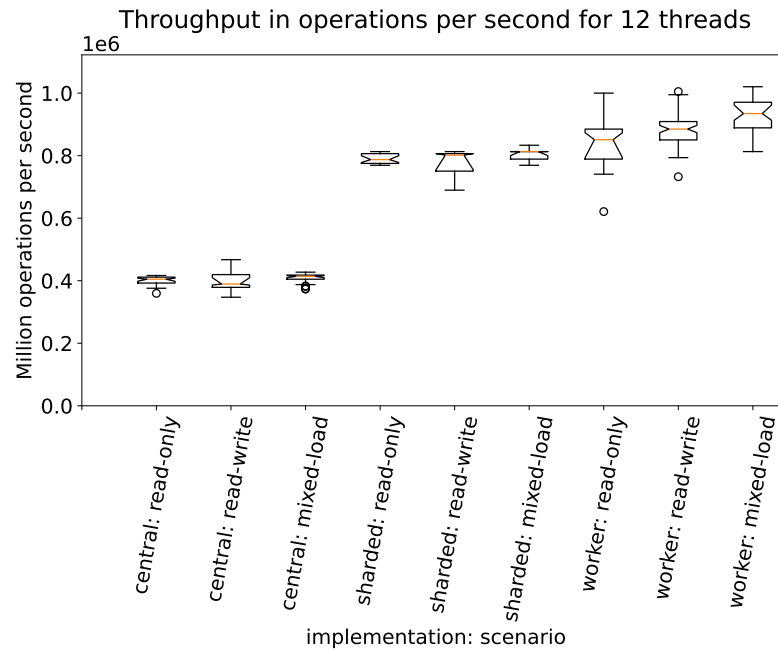

Figure 5: Comparing throughput of implementations on Omen

On Omen we can clearly see an improving trend in terms of the throughput when improving the implementations, however the returns seem to be diminishing. The difference between the central implementation and the sharded implementation is clear. The sharded implementation is a clear imrovement. Comparing the sharded and the worker implementation might require you to squint your eyes a little as the there is not always an improvement. Most of the the time there is an improvement but in some cases outliers perform worse than the sharded implementation. In the case of read only operations the throughput even overlaps with the sharded implementation.

3.4 Throughput per second depending on scheduler threads

For this experiment we use the same parameters mentioned in subsection 3.1 but set the scenario to read-write, as it provides the results of a balanced mix of operations and requires more operations to complete. We will also only look at the worker implementation as it seems to have the most potential for the highest throughput. In this experiment we measure the throughput of operations depending on the amount of scheduler threads. For Omen, we will look at 1-12 threads and for firefly 1-128.

Throughput per second of scenario: read-write, implementation: worker

Figure 6: Throughput Omen read-write

Looking at the results from Omen, we can see that there is a general trend that improves the throughput as we add more threads. However the returns are diminishing as the curve is logarithmic. In the last two measurements (11 and 12 threads) we can see that they are very similar. While having 12 threads makes it possible to outperform 11 threads, it is not implied as some results seem to overlap or are worse.

3.5 Speedup depending on scheduler threads

For this experiment we will use the same parameters mentioned in subsection 3.1 and similar variables as in subsection 3.4. We will only be looking at the worker implementation and at the read-write scenario.

Speedup of scenario: read-write, implementation: worker

4 Insight

4.1 Lightweight processes and its impact

The lightweight-ness of the processes has definately influenced my solution, because in other languages they have a hefty additional cost to create. This means that there is usually an optimal amount of processes you can have before your speedup starts to actually decrease. In Erlang however, the worker implementation can scale to 100 users and 100 buckets, having 201 processes in total without any issues and still packs a punch. Consequently I don't restrain myself as much as in the other languages in terms of the amount of processes created. If the processes did cost much higher, I would probably be looking for an optimal amount of processes to have to reach the optimal performance and base my implementation on that instead of going nuts with the amount of processes.

4.2 Distribution of cores over multiple machines

If the application would be distributed over multiple machines, the achitecture or the worker implementation would still be viable, but on a larger scale. This means that perhaps it would be better to give each process or each X amount of processes a machine. Since my implementation now primarily focusses on processing the most amount of requests in the least amount of time, I would probably shift my focus to avoiding network requests since these are responsible for most of the latency in operations. I would probaby be looking to keep a cache of the Bucket as close to the user as possible, Maybe even the bucket and the worker process as close together as possible, so I would probably be looking at minimizing physical distances as well. Some other techniques might be to replicate the buckets and have a load balancer in front of them.

To distribute processes to different machines, I would write a script that activates the Erlang shell with the right parameters. This allows your local computer to spawn processes on the remote

9

computer with an Erlang program, which is run after activating the Erlang shell on the remote machine.